





# When Metrics Say “Yes” But Humans Say “No”

---

A Case Study of Hallucination Reduction for Factual  
Summarization

Author: Ryan Powers

# Project Characters



BART(Generator)

Fine-tuned  
summarization  
model

FactCC(Verifier)

Classifier that scores  
article summary  
factual consistency.


CNN/Daily Mail  
(Dataset)

News articles with  
highlight summaries

Single Human  
Annotator

Single rater doing a  
50-example audit.

# Task/Problem/Goal



*Test whether better automatic metrics  $\Rightarrow$  better human-judged summaries.*

**Success Criteria:**  $\Delta \text{FactCC} \geq +2.0$ ,  $\text{ROUGE-L drop} \leq 1.0$

## Task:

CNN/DailyMail abstractive summarization (BART-base)

## Problem:

Summaries are fluent but sometimes factually wrong (hallucinations)


## Goal:

Improve factual consistency **without** retraining a bigger model, use reranking instead.

---

# Architecture Diagram (Generate → Verify)

- Generate  $K=5$  candidates with BART (beam search)
- Score each candidate with:
  - FactCC: probability the summary is consistent with the article.
  - RoBERTa-MNLI: entailment score (for ablation).
- Pick the candidate with highest FactCC score.



# BART + FactCC / Experimental Setup

## Models

- **Generator:** facebook/bart-base (fine-tuned on 20k CNN/DM, 3 epochs, lr=5e-5) -
  - a. **Verifier A:** FactCC (factual consistency) -
  - b. **Verifier B:** RoBERTa-MNLI (ablation)

## Data


- **Train:** 20k subset (seed=42) - **Val:** 2k | **Test:** 11,490

## Decoding

- **Beam search, K=5, max length 140**


**Success Criteria** ✓  $\Delta\text{FactCC} \geq +2.0$  points ✓ ROUGE-L drop  $\leq 1.0$  point

After analyzing K-ablation shows strong gains up to ~3–5, with diminishing returns.



# Automatic Results


- **ROUGE-L:**
  - LEAD-3: 24.9
  - BART baseline: 28.1
  - FactCC rerank: 27.9
  - NLI rerank: 28.1
- **FactCC (avg factuality score):**
  - Baseline: 0.427
  - FactCC rerank: 0.661 (+55%,  $p < 0.00001$ )
- **Success criteria satisfied:**
  - $\Delta\text{FactCC} \geq +2.0$  ✓
  - ROUGE-L drop  $\leq 1.0$  (-0.17) ✓



# Setting / $N = 50$ human audit

```
. df[df['changed']==True].sample(30) and  
df[df['changed']==False].sample(20)
```

For this project, a hallucination is any statement in the summary that is not supported by, or contradicts, the source article, including entity, event, temporal, and out-of-thin-air errors.



- 50 test examples:
  - 30 where baseline vs rerank disagree
  - 20 where they agree
- Human (Single Annotator) preference:
  - Tie: 27 (54%)
  - Baseline preferred: 14 (28%)
  - Reranked preferred: 9 (18%)
- Net hallucination effect: reranker fixes 1 error, introduces 3  $\rightarrow$  **Net -2**

# When Reranking Helps vs Hurts

## The Success (Logic Repair)

- **Context:** Refugees formed a human chain to stop attackers.
- **Baseline (The Mistake):** "15 men arrested... But were *prevented* [from being arrested] because they huddled together..."
  - Status: **Logical Error** (Implies the chain stopped the police).
- **Reranked (The Fix):** "...claimed the men tried to throw other Christians off... but were *prevented* [from doing so] because they huddled together..."
  - Status: **Corrected** (The verifier caught the causal mismatch).

## The Failure (Hallucination)

- **Context:** A father hit his toddler; store owner gave video to police.
- **Baseline (The Safe Bet):** "Justin Whittington... was taken into custody... relationship not yet revealed..."
  - Status: **Boring but Factual** (Stuck to the text).
- **Reranked (The Lie):** "...Police were notified... and asked the public for help identifying the man."
  - Status: **Hallucination** (The article never mentioned a public appeal).


---

Verifier (FactCC) isn't actually checking facts it's just checking for patterns.



# Conclusion

- **Lead bias:** CNN/DailyMail favors early sentences; LEAD-3 gets ROUGE-L 24.91 (only 3.18 points below BART). Reranking often moves toward extractive “lead-like” summaries without improving factuality.



## **FactCC mis-scores paraphrases:**

Penalizes legitimate paraphrasing, prefers extractive but slightly wrong spans.

**Limited beam diversity:** in ~32% of audited examples all K=5 candidates had hallucinations → reranking cannot fix generator-level failures.

- **Why:** Goodhart's Law. Once I made FactCC the optimization target, the system learned to game FactCC rather than be more truthful.

# References & Github

- Lewis et al. 2020. “BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation.” *ACL*.
- Hermann et al. 2015; Nallapati et al. 2016. CNN/DailyMail summarization dataset.
- Kryscinski et al. 2020. “Evaluating the Factual Consistency of Abstractive Text Summarization.” (FactCC).
- Ji et al. 2023. “Survey of Hallucination in Natural Language Generation.”
- Zha et al. 2023. “How Well Do Factuality Metrics Align with Human Judgments?” (for your QAGS/FactCC alignment discussion).

<https://github.com/rPowDS/MIDS266-NLP-final-project>